

MARS: A Reconfigurable Multi-Task-Capable Accelerator Architecture for Ultra-Low-Bit Streaming Inference

Po-Chun Huang, Pao-Ann Hsiung, and Chun-Hsian Huang

Abstract—Edge FPGA platforms must serve multiple AI tasks under tight resource budgets. The surveyed PR-based CNN methods require a reconfiguration controller and per-task partial bitstreams with millisecond-scale switch latencies, while fixed-feature-extractor designs cannot adapt the feature extractor at runtime. Combining ultra-low-bit quantisation with parameter-efficient transfer learning (PEFT) on a streaming dataflow accelerator could resolve this, but we observe a pronounced degradation of single-branch adapter effectiveness specifically at 1 bit. This paper presents MARS (Multi-Branch Adapter Reconfigurable Streaming-accelerator), an ultra-low-bit FPGA accelerator that integrates a custom Adapter IP into a streaming-dataflow backbone. MARS uses multi-branch adapter scaling to recover part of the transfer accuracy lost under binarisation, a Residual Correction (RC) term—the adapter’s pre-existing down-projection bias, realised at zero additional arithmetic or datapath cost, with its storage counted in the adapter budget—to mitigate the pronounced 1-bit-specific degradation, and a runtime configuration interface that reprograms all task-dependent parameters over a single bitstream without any fabric reconfiguration. On a frozen 1-bit VGG-like backbone, multi-branch scaling raises transfer accuracy by +3.4 percentage points on average (up to +11.3) across twenty backbone–target pairs in the deployed-geometry sweep, and RC contributes up to +23.0 pp. On a PYNQ-Z2, the compactness build switches at runtime among five task-specific models with on-board accuracy within 0.69 pp of software and a 1.86 ms per-task configuration rewrite (≈ 25 KB) with verified post-switch correctness; a throughput-oriented build sustains 1,866 img/s at 2.214 W (842.7 img/s/W) on a protocol-matched 1W1A workload with deployment-representative, different adapter geometries (Vivado-estimated FPGA power vs. measured GPU power).

Index Terms—Parameter-efficient transfer learning, ultra-low-bit quantisation, FPGA accelerator, streaming dataflow, runtime task switching, residual correction.

I. INTRODUCTION

EDGE FPGA accelerators are increasingly asked to serve several AI tasks from one device: a drone alternates between navigation and object recognition; a smart camera switches between day-time and infra-red models. On FPGAs, the prevailing answer is partial reconfiguration (PR): each task is compiled into a partial bitstream, and a fabric-reconfiguration controller swaps hardware regions at run time [1], [2]. PR works, but it carries a fixed tax—a reconfiguration controller, a per-task bitstream repository that grows linearly with the task count, and measured switch latencies

from 5.2 ms (ZyCAP) up to 97–114 ms (PCAP) on PYNQ-class devices. The alternative family fixes a feature extractor in the fabric and swaps only back-end logic [3] or selects among variants resident on chip [4]; these avoid fabric rewrites but freeze the reachable task set at synthesis time.

A third path is suggested by parameter-efficient transfer learning (PEFT) [5], [6], [7]: keep one shared backbone frozen and adapt tasks through small inserted modules such as Conv-Adapter [8]. If the backbone is also quantised to the 1-bit weight, 1-bit activation (1W1A) regime [9], [10], [11] and mapped to a streaming dataflow architecture [12], [13], then a task switch reduces to rewriting a small set of parameters—no fabric reconfiguration at all. However, in the evaluated VGG-like configurations, a single binary adapter branch exhibits a substantially larger transfer-accuracy degradation at 1W1A than at 2–8 bits, echoing the precision-redundancy perspective of [14].

This paper presents MARS (Multi-Branch Adapter Reconfigurable Streaming-accelerator), which makes the ultra-low-bit adapter path practical and carries it through to an on-board-verified system. Our contributions are:

- 1) **A 1-bit-specific empirical finding and its hardware-efficient mitigation.** The down-convolution bias of Conv-Adapter, isolated here as a Residual Correction (RC) term, has a substantially larger measured effect at 1W1A: it contributes up to +23.0 pp on CIFAR-10→SVHN in the deployed geometry, whereas its effect at 2–8 bits is typically within ± 0.3 pp (at most 1.4 pp). RC is realised at zero additional arithmetic or datapath cost by folding it into the accumulator initial value, with its storage and configuration bytes counted in the adapter budget.
- 2) **Multi-branch adapter scaling.** Summing M independently-parameterised binary branches increases the effective capacity of the binary adapter path, raising 1-bit transfer accuracy by +3.36 pp on average (up to +11.33 pp) across twenty backbone–target pairs in the single-seed deployed-geometry sweep.
- 3) **Single-bitstream runtime task switching.** A 19-LUT AXI-Lite Function Switch Controller reprograms all task-dependent parameters (≈ 25 KB per task) at run time; on-chip cost is $O(1)$ in the number of same-backbone, same-input-size, same-output-head classification tasks.
- 4) **On-board verification with explicit evidence boundaries.** A compactness build switches at runtime

among five task-specific models on one PYNQ-Z2 bitstream, each within 0.69 pp of its software reference, with a 1.86 ± 0.04 ms configuration rewrite (post-switch correctness verified); a throughput build sustains 1,866 img/s at 2.214 W on a matched 1W1A workload. Multi-branch configurations beyond $M=1$ exceed the XC7Z020 and are verified at the RTL-simulation and post-synthesis levels only.

The remainder of this paper is organised as follows. Section II reviews the three research threads MARS builds on. Section III presents the method: the binary adapter, RC, multi-branch scaling, and training. Section IV details the hardware architecture and the switching mechanism. Section V reports software, resource, on-board, and cross-platform results with explicit evidence levels, Section VI states the limitations, and Section VII concludes.

II. RELATED WORK

A. Ultra-Low-Bit FPGA Dataflow Accelerators

Binarised neural networks replace multiply–accumulate with XNOR–popcount [9], [10], and streaming-dataflow compilers such as FINN [12], [13] map every layer to a dedicated matrix–vector–activation unit (MVAU) with on-chip weights, yielding high throughput per watt on small FPGAs [15], [16]. Beyond FINN, RTL-level MVAU implementations [17], HPC-oriented extensions [18], hls4ml [19], ONNX-to-hardware flows [20], and all-on-chip stream accelerators [21], [22] broaden the dataflow family, and recent surveys catalogue the design space [23], [24], [25], [26], [27], [28], [29]. On the training side, the line from BinaryConnect [30] through Bi-Real Net and ReActNet [31], [32] and FPGA BNN systems such as FP-BNN [33] improves binary representational capability, while integer-quantisation practice is consolidated in [34], [35], [36], [37]. These pipelines, however, are compiled for a single task: weights are baked into the bitstream, and serving another task conventionally means another bitstream.

B. PEFT and Adapter Methods

Adapters [5], LoRA [6], and their vision descendants [7] adapt a frozen backbone with a small number of trained parameters. Conv-Adapter [8] inserts a down-projection/nonlinearity/up-projection bottleneck beside convolutional layers. Edge-oriented PEFT variants for CNNs and IoT devices are emerging [38], [39], [40], and benchmarks systematise the method space [41]. Prior PEFT work targets full-precision or moderately quantised models; the 1-bit regime—where the adapter itself must be binary—is largely unexplored, and [14] suggests the precision redundancy available to adapters becomes severely limited at 1 bit.

C. Task Switching on FPGAs

Published task-switching FPGA inference falls into three families: PR-based systems (ACNNE [1], coarse-to-fine PR [2], dynamic PR managers [42], [43]), fixed-feature-extractor designs with per-task back-ends [3], and on-chip

parameter selection among synthesis-time variants [4]. To the best of our knowledge among the surveyed FPGA/PEFT accelerator works, no prior system combines an ultra-low-bit shared backbone, feature-extractor-level adapter adaptation, and single-bitstream runtime task switching verified on hardware; MARS fills this gap.

III. THE MARS METHOD

A. Background: Streaming 1W1A Pipelines

In a FINN-style pipeline every layer is a matrix–vector–activation unit: an XNOR array against on-chip 1-bit weights, a popcount, and a thresholding stage that folds BatchNorm and the binary activation into one per-channel integer comparison. Critically, thresholding happens *inside* the MVAU, so the raw partial sum is not normally exposed; any additive per-task correction applied after an MVAU would have to re-quantise an already-binarised signal. This is the architectural obstacle that the Super Wrapper of Section IV removes, and it is why the adapter must be co-designed with the pipeline rather than appended to it.

B. One-Branch Binary Adapter

MARS inserts one adapter beside each of the five adapter-hosting convolution layers of a frozen 1W1A VGG-like backbone (the FINN reference CNV network: six convolutional and three fully connected layers, pre-trained on CIFAR-10). For hidden channel j the down-projection accumulates XNOR products of the binary input x with 1-bit weights and adds a per-channel bias:

$$h_j = b_j + \sum_{i=1}^{C_{in}} \text{XNOR}(W_{j,i}^{\text{down}}, x_i), \quad j = 1, \dots, C_h, \quad (1)$$

where $\text{XNOR}(\cdot, \cdot)$ denotes the bitwise exclusive-NOR and C_h is the bottleneck width. A binary activation $a_j = \mathbb{I}[h_j \geq 0]$ follows, and the up-projection restores the output width:

$$z_o = \sum_{j=1}^{C_h} \text{XNOR}(W_{o,j}^{\text{up}}, a_j), \quad o = 1, \dots, C_{out}. \quad (2)$$

The deployed geometry uses 1×1 down-convolutions with $C_h = C_{in}/4$ and a scalar per-branch scale.

C. Residual Correction (RC)

The bias term b_j in (1) is not new—it is the bias Conv-Adapter already has. Our finding is that at 1 bit this term changes character: without it, the binary activation can only cut the popcount distribution at zero, which need not coincide with the decision boundary the task requires; with it, the cut can move toward that boundary—consistent with the ablation, though not itself a verified mechanism. We therefore isolate it as the Residual Correction and keep it explicitly. A small example makes the mechanism concrete. Suppose a hidden channel’s popcount over the window ranges over $[0, 64]$ and the task-optimal decision boundary sits at 41. Without a bias, $\text{sign}(h)$ can only cut at zero of the *centred* accumulation—equivalent to a fixed boundary at 32—so every sample between




Fig. 1. RC in hardware: accumulator initialisation from `ram_rc`; no extra adder or pipeline stage.

32 and 41 lands on the wrong side, and at 1 bit there is no downstream precision to compensate. With $RC = -9$ folded into the accumulation, the cut moves onto the boundary—an illustrative numerical example rather than values extracted from a measured network channel. One plausible explanation for the smaller effect at ≥ 2 bits is that additional activation levels preserve magnitude information that downstream layers may exploit; this interpretation is consistent with the ablation but is not directly verified by activation-distribution analysis.

In hardware, RC is folded into the *initial value* of the down-projection accumulator: the accumulator must load an initial value at every output window anyway, so RC replaces the constant zero at zero additional arithmetic, DSP, or pipeline cost; the only real cost is the RC storage and its share of the per-task configuration bytes, which we count in the adapter budget. Fig. 1 shows the realisation: `ram_rc` drives the initialisation input of the multiplexer in front of the accumulator register, selected by the FSM at every window start, while the remaining cycles follow the accumulate loop—the crossed-out “+RC” stage does not exist in the RTL.

D. Multi-Branch Adapter

MARS instantiates M structurally identical branches with independent weights, RC, and scales, and sums them:

$$y_o = \sum_{i=1}^M \alpha_i z_o^{(i)}, \quad o = 1, \dots, C_{\text{out}}. \quad (3)$$

Summing independently parameterised branch outputs enriches the set of integer-valued corrections the adapter can express (each branch output is itself a multi-level popcount value, so no fixed output-level count is claimed). The parameter overhead is linear, $\Delta P(M) = M \times 58,533$ (+3.78% of the 1,546,690-parameter backbone per branch).

E. Wide vs. Deployed Geometry

Two adapter geometries are used deliberately. The *wide* geometry (3×3 down-convolution, per-channel α , $C_h = C_{\text{out}}/4$) is the software-oriented reference and achieves the highest accuracy on the CIFAR-10 backbone; the *deployed* geometry (1×1 down-convolution, scalar α , $C_h = C_{\text{in}}/4$) is the configuration carried into the bitstream, chosen because a 1×1 projection needs no additional line buffering in a streaming pipeline and a scalar α collapses into a single per-layer look-up table. The relative ranking of the two geometries is source-dependent (Section V), so conclusions involving exact accuracy differences are reported separately for each geometry.

F. Training

All task-dependent parameters—the M branches and the classifier—are trained jointly, end-to-end, with the backbone frozen. Branch weights are binarised in the forward pass,

$$W^{\text{down},(i)} = \text{sign}(\tilde{W}^{\text{down},(i)}), \quad W^{\text{up},(i)} = \text{sign}(\tilde{W}^{\text{up},(i)}), \quad (4)$$

and the backward pass uses the straight-through estimator

$$\frac{\partial \text{sign}(u)}{\partial u} \approx \mathbf{1}_{|u| \leq 1}, \quad (5)$$

so ADAM updates the real-valued latent weights and the real-valued RC bias. The per-branch scale α_i multiplies its branch output linearly and receives the exact gradient

$$\frac{\partial L}{\partial \alpha_i} = \left\langle \frac{\partial L}{\partial y}, z^{(i)} \right\rangle, \quad (6)$$

with no estimator involved. Nothing forces the branches apart: they share one structure and diverge only through independent random initialisation and joint training. At export, α_i is quantised into a Q8 contribution look-up table and the RC bias into the accumulator initial value.

G. Computational Cost

Each adapter branch adds 1×1 down- and up-projections across the five hosted layers. Two counts must be distinguished: the compute-then-crop *software reference* evaluates both projections over the full input feature map, totalling 4.23 MMACs per branch against the backbone’s 59.46 MMACs, whereas the deployed RTL’s centre-selection logic evaluates only the output positions, executing 3.01 MMACs per branch. With 1 MAC = 2 OPs, the unified per-image cost—using the software-reference count as a common normaliser—is

$$\text{OP}_{\text{total}} = 2 \times (59.46 + 4.23 M) \text{ MOPs}, \quad (7)$$

i.e., 127.38 MOPs at $M=1$ —a +7.1% compute overhead per branch, all of it XNOR–popcount.

IV. HARDWARE ARCHITECTURE

A. System Overview

Fig. 2 shows the deployed system on a Zynq XC7Z020 at 100 MHz. The streaming pipeline consists of the first-layer unit MVAU0, five adapter-enhanced Super Wrappers (Super1–5) hosting the Conv2–Conv6 layers, and the fully connected units MVAU6–MVAU8, connected by FINN-style streaming interfaces with input/output DMA. A Function Switch Controller (FSC) on the AXI-Lite bus fans a configuration write interface out to all task-dependent memories.

B. Adapter-Enhanced Super Wrapper

Fig. 3 shows the Super Wrapper. A splitter duplicates the input stream to (i) a modified FINN MVAU core whose in-module thresholding is decoupled so that the raw 32-bit partial sum is exposed, and (ii) the Adapter IP, whose 8-bit contribution passes through a small synchronisation FIFO. A stream adder fuses the two contributions and applies the Q8




Fig. 2. MARS streaming dataflow on the Zynq PL: MVAU0, adapter-enhanced Super1–5, and MVAU6–8, with the Function Switch Controller driving the configuration bus.




Fig. 3. MVAU + Adapter Super Wrapper: the modified MVAU core exposes raw partial sums; the adapter contribution is fused before the shared Q8 threshold.

threshold, producing the 1-bit output; binarisation therefore happens *after* the adapter contribution is added, which is what makes a post-accumulator adapter possible in a streaming binarised pipeline. The adapter’s weight, RC, and α -LUT memories are configuration-writable distributed RAM. Denoting the backbone partial sum by p_o and the up-projection popcount by q_o , the fused output bit is

$$\hat{y}_o = \mathbb{I}[(p_o \ll 8) + s_o \cdot \text{LUT}_\alpha(q_o) \geq T_o^{Q8}], \quad (8)$$

where LUT_α is the per-layer Q8 contribution table that absorbs α , s_o a per-channel sign, and T_o^{Q8} the BatchNorm-fused threshold in Q8. Everything on the right-hand side is either a shift, a table look-up, or an integer add—no multiplier is instantiated.

Fig. 4 details the Adapter IP datapath: a two-stage pipeline in which the down-projection XNOR–popcount accumulates onto an accumulator whose initial value is the RC bias, the sign stage produces the binary hidden vector, and the up-projection XNOR–popcount feeds a Q8 α -LUT that emits the branch contribution. All weight, RC, and LUT memories are configuration-writable distributed RAM fed by the FSC.

C. Function Switch Controller

The FSC is an AXI-Lite slave with a 64 KB aperture. A write FSM latches the 16-bit address, decomposes it combinatorially into a 3-bit unit select and an 11-bit word address, and issues a single-cycle write-enable pulse to one of nine destinations (MVAU0 thresholds, MVAU8 classifier weights, the five Super adapters, and the MVAU6/MVAU7 thresholds) while broadcasting the shared address/data bus. The entire




Fig. 4. Adapter IP internal datapath. The RC bias enters as the accumulator initial value; the α -LUT converts the up-projection popcount into the Q8 branch contribution.

TABLE I
FSC ADDRESS MAP (64 KB APERTURE).

Unit	Byte Range	Target
0 low	0x0000--0x0FFF	MVAU0 thresholds (64 w)
0 high	0x1000--0x1FFF	MVAU8 classifier weights (1024 w)
1–5	0x2000--0xBFFF	Super1–5 adapter: enable, RC, down/up weights, thresh, sign, α -LUT
6	0xC000--0xDFFF	MVAU6 thresholds (512 w)
7	0xE000--0xFFFF	MVAU7 thresholds (512 w)

switching infrastructure costs 19 LUTs, zero BRAM, and zero DSP, and is orthogonal to the streaming datapath: no inference-path logic is touched.

Fig. 5 shows the internal circuit. Table I gives the address map.

D. Runtime Task Switching

Because every eligible task shares the same frozen backbone, the binary convolution and FC weights baked into the bitstream are identical across tasks; only the BatchNorm-fused thresholds, the classifier projection, and the adapter parameters differ. A full task switch is therefore a ≈ 25 KB configuration rewrite (25,088 stored bytes; 6,757 32-bit word writes, i.e. 27,028 bus bytes, since sub-word fields are widened to full words on the configuration bus), shown in Fig. 6: wait for the output-DMA done flag, bulk-write the new task blob over




Fig. 5. Function Switch Controller: AW latch, write FSM, combinational address slicing, and a 1-of-9 write-enable demultiplexer over a shared configuration bus.

TABLE II
PER-LAYER PE FOLDING OF THE TWO BUILDS.

Unit	Throughput PE	Compactness PE
MVAU0	16	1
Super1	32	1
Super2	16	1
Super3	16	1
Super4	4	1
Super5	1	1
MVAU6	1	1
MVAU7	1	1
MVAU8	5	5

TABLE III
ADAPTER MEMORIES (EXAMPLE: SUPER4, PE=4).

RAM	Depth	Width	Contents
ram_down	128	32 b	down-projection weights
ram_up	64	128 b	packed up-projection weights
ram_rc	32	16 b	RC bias (signed)
contrib_lut	256	32 b	Q8 contribution (α absorbed)

TABLE IV
TWO-TASK VS. N -TASK SWITCHING: SAME MECHANISM, TWO DEPTHS.

	Throughput (2-task)	Compactness (N -task)
Folding	heterog. PE ≤ 32	uniform PE = 1
Baked at synthesis	one adapter task	nothing (memories empty)
Moved per switch	thr. + classifier ≈ 15 KB	full blob ≈ 25 KB
Reachable models	exactly two	any eligible same-backbone task
Verified	on-board, 2 tasks	on-board, 5 tasks

MMIO, and launch the next batch—no additional flush beyond the output-DMA-done barrier and no fabric reconfiguration. Each additional task costs only an off-line host-side blob; the on-chip cost is $O(1)$ in the task count, bounded by the shared frozen backbone, the $32 \times 32 \times 3$ input format, the shared 10-class output-head interface, and the classification-task assumption.

Two builds of the same architecture are used consistently: a *throughput* build with FINN-default heterogeneous folding (PE up to 32), and a *compactness* build with uniform PE=1 convolutions that trades throughput for the slice headroom hosting fully runtime-writable per-task state. Table IV summarises how the two builds realise two depths of the same mechanism. Table II lists the per-layer folding (visualised in Fig. 7), and Table III the adapter memories of one representative Super Wrapper; all adapter state is distributed RAM, which is why BRAM does not grow with the adapter.

E. Design Choices

Three choices deserve justification. *Distributed RAM over BRAM*: adapter memories are small, wide, and read every cycle; LUT-RAM keeps them adjacent to their datapaths and leaves the BRAM budget to the backbone memstreams, which is why the adapter adds zero BRAM. *Splitter + FIFO*

synchronisation: the adapter side-stream must rejoin the folded MVAU stream cycle-accurately under backpressure; a small synchronisation FIFO absorbs the latency skew so that the fused threshold sees matched operands, which is what keeps adapter-enabled throughput within $\approx 0.5\%$ of the backbone. *Q8 fixed point*: the selected eight-fractional-bit format keeps the measured software-to-board gap below one percentage point in the deployed experiments (Table XII); a systematic precision sweep of this scaling was not performed.

V. EXPERIMENTAL RESULTS

A. Setup

Software experiments use PyTorch with Brevitas [44] quantisation-aware training (200 epochs, ADAM, lr 0.005, step decay, seed 2024); hardware is generated through FINN and Vivado 2022.2 and deployed on a PYNQ-Z2 (XC7Z020). Five $32 \times 32 \times 3$ classification datasets are used: CIFAR-10 (backbone source), SVHN, FashionMNIST, STL10, and CINIC10. Functional correctness is established at three levels: per-module RTL simulation (43,520 vectors, 100% pass), top-level dataflow simulation against the PyTorch golden model, and end-to-end on-board execution.




Fig. 6. End-to-end runtime task-switching sequence.




Fig. 7. Folding difference between the builds: PE=1 (compactness) vs. PE-parallel lanes (throughput, illustrated for PE=4); a unit completes PE×SIMD binary operations per cycle.

TABLE V
DATASETS (ALL ADAPTED TO A $32 \times 32 \times 3$ INPUT).

Dataset	Classes	Train / Test	Notes
CIFAR-10	10	50,000 / 10,000	natural objects (backbone)
SVHN	10	73,257 / 26,032	house-number digits
STL10	10	5,000 / 8,000	ImageNet-derived, 96^2 native
FashionMNIST	10	60,000 / 10,000	clothing, grayscale
CINIC10	10	90,000 / 90,000	CIFAR + ImageNet mix

Table V lists the five datasets; all are adapted to the backbone’s $32 \times 32 \times 3$ input (STL10 downscaled, FashionMNIST resized and channel-replicated). For each transfer, the backbone is frozen and the adapter and classifier are trained for 200 epochs with SqrHinge loss. Following the original BNN-PYNQ trainer flow, the best epoch is selected by test-set accuracy (no held-out validation split), so absolute accuracies are best-of-200-epoch operating points and may be optimistically biased; every configuration uses the identical rule, making comparisons procedurally consistent, although best-epoch selection on the test split may also bias relative differences.

B. Three-Level Functional Verification

Correctness is established bottom-up. (i) *Module level*: each Super Wrapper is simulated standalone against a memblock-faithful golden model—43,520 output vectors across the five adapter-hosting layers, 100% bit-exact. (ii) *Top level*: the complete stitched pipeline (DMA, MVAU0, Super1–5, pooling, MVAU6–8, label select) is simulated as one block design against the PyTorch reference. (iii) *Board level*: the generated bitstream executes the full test sets end-to-end. One methodological note: when RTL sources inside a stitched IP change, Vivado’s per-IP out-of-context synthesis cache must be reset explicitly, otherwise a stale netlist silently masks configuration writes—a failure mode worth documenting for reproducibility.

TABLE VI
RC ABLATION AT 1 BIT, WIDE GEOMETRY (CIFAR-10 $\rightarrow$ SVHN).

Adapter	M	No-RC	With RC	Δ_{RC}	Full-FT
Conv-Adapter [8]	1	63.06%	79.23%	+16.17	
	2	66.12%	83.01%	+16.89	94.91%
MARS	3	67.79%	83.48%	+15.69	
	4	69.08%	84.44%	+15.36	

TABLE VII
MULTI-BRANCH SCALING ON THE CIFAR-10 BACKBONE (WIDE GEOMETRY, SOFTWARE, 1W1A, RC).

Target	[8]		MARS				Full-FT
	$M=1$	$M=2$	$M=3$	$M=4$			
SVHN	79.23	83.01	83.48	84.44			94.91
STL10	67.50	67.48	67.68	68.15			71.05
FashionMNIST	82.71	83.57	83.63	83.60			92.36
CINIC10	64.50	64.71	64.85	65.02			68.02

C. Residual Correction at 1 Bit

Table VI isolates RC on CIFAR-10 $\rightarrow$ SVHN in the wide software geometry. RC contributes +15 to +17 pp at every branch count, and +23.0 pp in the deployed (1×1 , scalar- α) geometry; at ≥ 2 -bit precision the same term shifts accuracy typically within ± 0.3 pp (maximum observed 1.4 pp), demonstrating a substantially larger effect at 1W1A in the evaluated sweep. On the other transfer targets, RC gains range from +0.6 to +7.4 pp.

D. Multi-Branch Scaling in Software

Table VII reports the wide-geometry software branch sweep on the CIFAR-10 backbone, with the largest gains on hard transfers and diminishing returns as M grows. The headline numbers come from the *deployed-geometry* sweep over all twenty backbone–target pairs spanning five source backbones (Table XI): multi-branch scaling improves 1-bit transfer accuracy by +3.36 pp on average over the single-branch Conv-Adapter baseline, up to +11.33 pp (FashionMNIST $\rightarrow$ SVHN).

Table VIII extends the wide-geometry software study to the remaining sixteen pairs across the four additional source backbones, with gains up to +7.63 pp (STL10 $\rightarrow$ SVHN). In the deployed-geometry sweep of Table XI, every one of the twenty pairs improves, up to +11.33 pp on FashionMNIST $\rightarrow$ SVHN. Five-seed paired checks on two deployment-relevant transfers provide preliminary evidence of a positive $M=1 \rightarrow 4$ gain under the same test-selected checkpoint protocol: +6.64 $\pm$

0.43 pp on SVHN (95% CI $[+6.10, +7.18]$) and $+2.69 \pm 0.36$ pp on FashionMNIST (95% CI $[+2.25, +3.13]$); a selection-free final-epoch reading of the same training logs yields consistent paired gains ($+6.87 \pm 0.50$ and $+2.99 \pm 0.40$ pp). On the CIFAR-10 deployment backbone the deployed (1×1 , scalar- α) geometry preserves the broad target-wise trend of the wide geometry; elsewhere the ranking is source-dependent, and under severe mismatch the deployed geometry can exceed the wide one. All hardware results below use the deployed geometry.

E. Resource, Power, and Branch Cost

Table IX compares four post-implementation builds on the XC7Z020. Integrating the adapter and the switching infrastructure raises the throughput build to 78.4% LUTs at 99.8% slice occupancy; the deployed compactness build reaches 75.1% LUTs while providing the runtime-writable state for one active task; the five task images reside on the host and are loaded in turn. All builds use zero DSPs, and the adapter does not increase BRAM: all adapter memories live in distributed RAM. On-chip power (Vivado post-route estimate) is 2.215 W for the adapter-enabled throughput build ($+0.101$ W over its backbone) and 1.802 W for the compactness build. Post-route hierarchical analysis attributes roughly 25,593 LUTs ($\approx 48\%$ of the chip) to the five adapter-hosting Super Wrappers, dominated by the adapter weight RAMs, XNOR-popcount trees, and threshold comparators, consistent with MVAU resource-prediction studies [45]; occupied-slice utilisation (97–100%) saturates before LUT utilisation because control-set constraints spread the design across the placement grid. Table X quantifies the linear per-branch cost; $M \geq 2$ exceeds the XC7Z020 post-synthesis and is verified at the RTL-simulation and post-synthesis levels only (not placed-and-routed or on-board), and Fig. 8 plots the resulting accuracy-versus-LUT trade-off.

On the throughput build, integrating the adapter and the switching infrastructure raises total on-chip power by only $+0.101$ W ($+4.8\%$, $2.114 \rightarrow 2.215$ W), the dynamic term rising from 1.949 to 2.046 W; the FSC itself adds negligible power. (The 2.215 W build includes the FSC; the cross-platform comparison of Table XVI uses the 2.214 W adapter-enabled throughput build without it.) The compactness build sits lower (1.78–1.80 W) because uniform $PE=1$ folding instantiates far less parallel logic. The $\approx 17.4\times$ throughput difference between the builds ($107 \rightarrow 1,859$ img/s on the backbone) is the price of the folding design point, not an optimisation introduced here; the pipelining contribution of this work is that adapter-enabled throughput stays within $\approx 0.5\%$ of the corresponding backbone, because the adapter side-stream is cycle-accurately synchronised into the folded datapath.

Fig. 9 shows the post-implementation placement of the compactness build: adapter logic and backbone MVAUs dominate and interleave across the die, while the Function Switch Controller appears only as a negligible scattering of cells, matching its 19-LUT footprint.




Fig. 8. Accuracy gain vs. post-synthesis LUT demand for multi-branch scaling (compactness build). Large-gain: $\Delta_{\text{multi}} \geq 3$ pp (9 pairs); the remaining 11 pairs form the small-gain group. $M \geq 2$ exceeds the XC7Z020.




Fig. 9. Post-implementation layout of the MARS compactness build on XC7Z020. Yellow: adapter datapaths and adapter RAMs; green: MVAUs; purple: Function Switch Controller; orange: PS; blue: other SoC-integration logic.

F. On-Board Runtime Task Switching

Table XII reports on-board accuracy for the five task-specific models served from the single compactness bitstream (loaded in turn from host-resident images), each measured over its 10,000-image test set (STL10 over its full 8,000). Every task reproduces its software reference within 0.69 pp; the residual gap is consistent with accumulated fixed-point and threshold-quantisation effects rather than an observed RTL mismatch. Each accuracy figure is the mean of ten runs at batch 1000 over the full test set; the on-board digits are stable across runs. Per-module RTL simulation is bit-exact, so no RTL mismatch was observed; the residual gaps are consistent with Q8 threshold rounding reclassifying a few borderline images. Fifty random task switches yield a configuration-rewrite time of 1.86 ± 0.04 ms (post-switch inference correctness verified separately), dominated by the single-beat AXI-Lite write path. Table XIII positions MARS against published task-switching FPGA inference: MARS is the only surveyed system that

TABLE VIII
MULTI-BRANCH ADAPTER SCALING ACROSS OTHER SOURCE BACKBONES (WIDE GEOMETRY, 1W1A, RC).

Backbone	Target	Conv-Adapter [8]	MARS				Δ_{multi} (pp)	Full-FT
		$M=1$	$M=2$	$M=3$	$M=4$			
SVHN	CIFAR-10	46.34	49.54	49.56	50.41	+4.07	81.70	
SVHN	STL10	35.16	37.48	37.60	38.46	+3.30	63.14	
SVHN	FashionMNIST	81.47	82.40	82.86	83.02	+1.55	92.32	
SVHN	CINIC10	32.76	34.47	36.84	39.70	+6.94	67.85	
STL10	CIFAR-10	68.92	69.97	70.58	70.99	+2.07	83.01	
STL10	FashionMNIST	84.71	85.77	86.28	86.35	+1.64	92.73	
STL10	SVHN	73.22	77.53	80.16	80.85	+7.63	95.32	
STL10	CINIC10	56.07	56.74	57.39	57.36	+1.29	70.18	
FashionMNIST	CIFAR-10	30.42	30.54	32.64	33.26	+2.84	78.37	
FashionMNIST	STL10	30.84	31.94	33.20	31.81	+0.97	59.11	
FashionMNIST	SVHN	39.16	43.29	45.68	45.68	+6.52	94.36	
FashionMNIST	CINIC10	24.82	25.33	27.28	26.34	+1.52	64.63	
CINIC10	CIFAR-10	78.05	78.08	78.47	78.59	+0.54	80.92	
CINIC10	SVHN	67.43	70.60	71.12	72.30	+4.87	94.62	
CINIC10	STL10	66.42	66.86	67.33	67.83	+1.41	70.38	
CINIC10	FashionMNIST	79.81	80.89	81.26	81.59	+1.78	91.99	

TABLE IX
FPGA RESOURCE UTILISATION ACROSS FOUR MARS BUILD CONFIGURATIONS (POST-IMPLEMENTATION, XC7Z020).

Configuration	Function Switch	Occupied Slices	Slice LUTs	Slice Reg. (FF)	BRAM
Backbone (Throughput)	no	11,648 (87.58%)	25,322 (47.60%)	31,621 (29.72%)	99.5 (71.07%)
MARS (Throughput, 2-task)	yes	13,272 (99.79%)	41,729 (78.44%)	53,285 (50.08%)	99 (70.71%)
Backbone (Compactness)	no	8,926 (67.11%)	15,651 (29.42%)	22,219 (20.88%)	85 (60.71%)
MARS (Compactness, N-task)	yes	12,957 (97.42%)	39,967 (75.13%)	40,926 (38.46%)	84 (60.00%)

TABLE X
PER-BRANCH PARAMETER AND POST-SYNTHESIS COST (COMPACTNESS BUILD).

M	Params	vs. Backbone	Slice LUTs (%)	Power (W)
1	58,533	+3.78%	47,931 (90.1%)	1.841
2	117,066	+7.57%	71,841 (135.0%)	1.897
3	175,599	+11.35%	95,593 (179.7%)	1.915
4	234,132	+15.14%	125,538 (236.0%)	1.969

needs no fabric-PR controller while keeping the task set open after synthesis.

Table XIV reports the post-synthesis scaling of both builds for $M=1-4$: growth is linear and predictable in LUT/FF with BRAM constant and zero DSPs throughout, because every added branch lives entirely in logic and distributed RAM. Synthesis-stage figures overstate the final cost—after place-and-route the deployed $M=1$ builds settle at 78.4% (throughput) and 75.1% (compactness) of the device LUTs.

The 1.86 ms switch amortises to $\approx 0.02\%$ of a batch-1000

run, and the write path sustains one word per ≈ 275 ns. Storage also scales favourably: serving T tasks requires one full bitstream plus T host-side 25 KB blobs, versus T partial bitstreams (each orders of magnitude larger) for PR-based systems, and versus synthesis-time-fixed on-chip variants for parameter-selection designs.

G. Cross-Platform Comparison

Table XVI compares MARS against Conv-Adapter running the same 1W1A SVHN workload on an RTX 4090 and a Jetson Orin NX under one aligned protocol (matched input, batch 1000, end-to-end timing). The throughput build sustains 1,866 img/s at 2.214 W—842.7 img/s/W, $11.9\times$ the RTX 4090 and $5.9\times$ the Orin NX. FPGA power is the Vivado post-route estimate whereas GPU/Jetson power is runtime-measured (NVML/tegrastats), so the efficiency ratios are workload-specific and methodology-dependent; we claim them only for the measured 1W1A workload. Model geometry also differs by platform: the GPU/Jetson run the prior-art Conv-Adapter (3×3 wide geometry) while the FPGA runs

TABLE XI
DEPLOYED-GEOMETRY MULTI-BRANCH TRANSFER ACCURACY (%) OVER
ALL TWENTY BACKBONE-TARGET PAIRS ($M=1-4$, RC ON; SINGLE
SEED). $\Delta_{\text{MULTI}} = M=4 - M=1$ (pp).

Backbone	Target	$M=1$	$M=2$	$M=3$	$M=4$	Δ_{multi}
CIFAR-10	SVHN	73.72	77.16	78.77	79.81	+6.09
	FashionMNIST	80.42	82.00	82.82	83.54	+3.12
	STL10	67.46	67.58	68.00	68.06	+0.60
	CINIC10	64.80	65.16	65.38	65.45	+0.65
SVHN	CIFAR-10	43.18	45.13	46.93	48.37	+5.19
	FashionMNIST	78.51	79.76	80.56	81.89	+3.38
	STL10	34.41	34.60	35.45	35.10	+0.69
	CINIC10	33.59	35.21	34.29	38.03	+4.44
STL10	CIFAR-10	66.65	67.90	68.74	68.79	+2.14
	FashionMNIST	81.48	82.96	83.58	84.19	+2.71
	SVHN	68.15	73.18	74.17	75.65	+7.50
	CINIC10	54.26	55.33	56.17	56.56	+2.30
FashionMNIST	CIFAR-10	33.63	34.67	35.84	36.76	+3.13
	STL10	32.64	32.29	33.64	34.11	+1.47
	SVHN	51.37	58.19	59.42	62.70	+11.33
	CINIC10	27.75	28.37	29.55	29.56	+1.81
CINIC10	CIFAR-10	78.21	78.60	78.87	78.72	+0.51
	SVHN	69.14	73.88	75.39	76.34	+7.20
	STL10	67.21	67.60	67.54	67.42	+0.21
	FashionMNIST	80.23	81.86	82.17	82.88	+2.65
Mean gain over $M=1$ (pp)		—	+1.73	+2.52	+3.36	—

TABLE XII
COMPACTNESS BUILD: ON-BOARD N -TASK SWITCHING ACCURACY.

Dataset	Mode	Software Reference	On-Board	Gap (pp)
CIFAR-10	backbone only	81.46%	80.99%	-0.47
SVHN	MARS	73.72%	73.21%	-0.51
FashionMNIST	MARS	80.42%	79.98%	-0.44
STL10	MARS	67.46%	66.77%	-0.69
CINIC10	MARS	64.80%	64.64%	-0.16

the deployed 1×1 MARS geometry—a deployment-level comparison, not a same-model sweep. The two geometries also differ in software accuracy (79.23% vs. 73.72% best-epoch CIFAR-10→SVHN reference), so the efficiency ratios compare deployment-representative implementations rather than accuracy-matched models.

H. Deployment Guidance

The two builds are complementary operating points rather than competitors. If the workload is a single task—or a task pair known at synthesis time—the throughput build is the right choice: it delivers 1,866 img/s and the best energy efficiency in Table XVI. If the device must serve an open, growing set of same-backbone tasks, the compactness build is the only configuration whose *entire* task state is runtime-writable on the XC7Z020; its throughput cost is the folding, not the switching machinery. On larger devices the two roles may merge: the synthesis results suggest that a device offering $\approx 135\%$ of an XC7Z020’s LUTs could accommodate a throughput-folded design with fully writable adapter state, though this remains to be confirmed by place-and-route on the target platform. Regarding branch count, $M=1$ is the deployable point on this device; investing in $M>1$ pays off primarily on hard transfers (Table VIII), so branches should be added only where a software profile of the target transfer justifies them.

I. Fairness and Scope of the Comparison

Five conditions bound the cross-platform claim: (i) all platforms run the same dataset at the same 1W1A precision, each with its deployment-representative (and hence different) adapter geometry; (ii) the batch-1000 protocol and preprocessing are identical; (iii) timing is end-to-end, including host-side data movement; (iv) the power methodologies are asymmetric (Vivado estimate vs. NVML/tegrastats measurements); and (v) the energy-efficiency advantage is reported relative to the two measured GPU baselines and for this workload only—not as leadership over all XC7Z020-class devices or general superiority over GPUs.

J. Claim–Evidence Summary

Table XV consolidates what is experimentally established at each verification level, following the principle that every headline number should carry its evidence level and limitation.

VI. LIMITATIONS

The hardware-verified point is the $M=1$ compactness build; $M \geq 2$ configurations exceed the XC7Z020 and are verified at the RTL-simulation and post-synthesis levels only. The switch latency is not yet sub-millisecond: the 1.86 ms figure is bound by the single-beat AXI-Lite path, and a faster bulk driver was observed to corrupt adapter state, indicating that sub-millisecond switching requires an RTL change rather than a driver change. “Arbitrary” task counts refer to the resource dimension only—eligible tasks must share the frozen backbone, the $32 \times 32 \times 3$ input, the output-head interface, and the classification setting, host-side blob storage grows linearly, and achievable accuracy is bounded by the transfer capacity of the shared backbone. Reported accuracies are best-epoch operating points selected on the test split (the original BNN-PYNQ trainer flow); a validation-selected (or last-epoch) protocol is left for future work. Finally, FPGA power is a Vivado estimate rather than a board measurement.

VII. FUTURE WORK

Four directions follow directly. (i) A bare-metal AXI4-burst switching driver: the 1.86 ms figure is dominated by single-beat AXI-Lite writes, and a burst-capable configuration port will be evaluated as a path toward sub-millisecond switching and, ultimately, per-inference granularity. (ii) Multi-branch hardware verification on a larger device (e.g., Zynq Ultra-Scale+), turning the RTL/post-synthesis evidence for $M \geq 2$ into board evidence. (iii) Broader backbone support beyond the VGG-like CNV topology, including residual networks [46], which requires adapter placement across skip connections. (iv) Adapter-aware training that anticipates the Q8 export, closing the residual software-to-board gap below the current 0.69 pp.

VIII. CONCLUSION

MARS shows that ultra-low-bit streaming inference and parameter-efficient task adaptation can be combined on a single FPGA bitstream. A 1-bit-specific residual-correction term, realised at zero additional arithmetic or datapath cost,

TABLE XIII
TASK-SWITCHING FPGA INFERENCE: MARS VS. PUBLISHED APPROACHES.

Approach	Year	Switch Latency	Fabric-PR Controller	Task Storage Scaling	Adaptation Level
ACNNE [1]	2024	97–114 ms	Yes (PCAP)	partial bitstream per task	layer-block swap
CoarseToFine [2]	2024	5.2 ms	Yes (ZyCAP)	partial bitstream per task	model-level swap
Boudjadar et al. [42]	2025	PR-bound	Yes (DPR)	partial bitstream per task	model-level swap
Kang & Park [43]	2025	PR-bound	Yes (DPR)	partial-update masks per task	partial model update
FixyNN [3]	2019	N/A (fixed FE)	No	back-end hardware per task	fixed FE; per-task back-end
NeuroAdaptiveNet [4]	2025	0.23–0.67 μ s	No	all variants on-chip (fixed at synthesis)	model/variant select
MARS	2026	1.86 ms	No	$O(1)$ on-chip; ≈ 25 KB host blob/task	FE-level adapters

TABLE XIV
POST-SYNTHESIS RESOURCE AND POWER SCALING OF MULTI-BRANCH MARS ON XC7Z020 (BRAM CONSTANT, ZERO DSP).

Build	M	Slice LUTs	Slice Reg. (FF)	BRAM	Power (W)
Throughput (2-task)	1	73,169 (137.5%)	65,717 (61.8%)	99	2.135
	2	109,731 (206.3%)	86,727 (81.5%)	99	2.193
	3	145,061 (272.7%)	107,471 (101.0%)	99	2.337
	4	187,521 (352.5%)	128,698 (121.0%)	99	2.468
Compactness (N-task)	1	47,931 (90.1%)	42,922 (40.3%)	84	1.841
	2	71,841 (135.0%)	60,390 (56.8%)	84	1.897
	3	95,593 (179.7%)	77,830 (73.2%)	84	1.915
	4	125,538 (236.0%)	95,285 (89.6%)	84	1.969

TABLE XV
CLAIM–EVIDENCE SUMMARY.

Claim	Evidence Level	Dataset / Workload	Limitation
RC contribution (+16.2 pp wide, +23.0 pp deployed, $M=1$)	software bias on/off ablation	CIFAR-10→SVHN, 1W1A; other targets +0.6–7.4 pp	single seed; specific to 1 bit
Multi-branch gain (mean +3.4 pp, up to +11.3 pp)	software, single seed over 20 pairs; 5-seed paired $M=1 \rightarrow 4$ checks on two transfers	20 backbone–target pairs	$M \geq 2$ not on-board verified
$M \geq 2$ hardware cost	RTL simulation (bit-exact) + post-synthesis	$M=2-4$, both builds	exceed the XC7Z020 capacity; no place-and-route
$M=1$ on-board correctness	on-board, full test sets $\times 10$ runs	five datasets, compactness build	within 0.69 pp of software (fixed-point/threshold quantisation)
Runtime task switching	on-board, 50 random switches	five task models, one bitstream, 1.86 $\pm$ 0.04 ms	not sub-ms; same-backbone
Throughput / energy efficiency	measured end-to-end, batch 1000	1W1A SVHN, MARS vs. Conv-Adapter	32 $\times$ 32 $\times$ 3 classification only FPGA power Vivado-estimated

TABLE XVI
DEPLOYMENT-LEVEL COMPARISON AT MATCHED DATASET, PRECISION, INPUT SIZE, AND BATCH PROTOCOL BUT DIFFERENT ADAPTER GEOMETRIES: MARS (FPGA) VS. CONV-ADAPTER (GPU/JETSON), SVHN, BATCH 1000.

Metric	RTX 4090	Orin NX	MARS (Compact.)	MARS (Thrput.)
Adapter geometry	3 $\times$ 3 wide	3 $\times$ 3 wide	1 $\times$ 1 depl.	1 $\times$ 1 depl.
FPS (img/s)	9,388	2,752	107.1	1,866
Time/img (ms)	0.107	0.363	9.338	0.536
Power (W)	132.48	19.23	1.802	2.214
Efficiency (img/s/W)	70.86	143.1	59.4	842.7

and multi-branch adapter scaling recover part of the accuracy that binarisation removes from single adapters; a 19-LUT Function Switch Controller turns task switching into a ≈ 25 KB parameter rewrite with no fabric reconfiguration. On a PYNQ-Z2, five task-specific models run from one compactness-build bitstream within 0.69 pp of software, with 1.86 ms configuration-rewrite switching (post-switch correctness verified); separately, the throughput build reaches 842.7 img/s/W on the matched 1W1A workload. Future work includes a burst-capable switching driver, multi-branch verification on larger FPGAs, and broader backbone support.

REFERENCES

- [1] C.-H. Huang, S.-W. Tang, and P.-A. Hsiung, "ACNNE: An Adaptive Convolution Engine for CNNs Acceleration Exploiting Partial Reconfiguration on FPGAs," in *Proceedings of the IEEE International Symposium on Circuits and Systems (ISCAS)*, May 2024, pp. 1–5.
- [2] A. Rangsikunpum, S. Amiri, and L. Ost, "A Reconfigurable Coarse-to-Fine Approach for the Execution of CNN Inference Models in Low-Power Edge Devices," *IET Comput. Digit. Tech.*, vol. 2024, no. 1, art. no. 6214436, Jan. 2024, doi: 10.1049/cdt2/6214436.
- [3] P. N. Whatmough et al., "FixyNN: Energy-Efficient Real-Time Mobile Computer Vision Hardware Acceleration via Transfer Learning," in *Proceedings of the Conference on Machine Learning and Systems (MLSys)*, Apr. 2019.
- [4] A. El Bouazzaoui, O. Mouhib, and A. Hadjoudja, "NeuroAdaptiveNet: A Reconfigurable FPGA-Based Neural Network System with Dynamic Model Selection," *Chips*, vol. 4, no. 2, art. no. 24, May 2025, doi: 10.3390/chips4020024.
- [5] N. Houlsby et al., "Parameter-Efficient Transfer Learning for NLP," in *Proceedings of the International Conference on Machine Learning (ICML)*, Jun. 2019, pp. 2790–2799.
- [6] E. J. Hu et al., "LoRA: Low-Rank Adaptation of Large Language Models," in *Proceedings of the International Conference on Learning Representations (ICLR)*, Apr. 2022.
- [7] B. X. B. Yu et al., "Visual Tuning," *ACM Comput. Surv.*, vol. 56, no. 12, art. no. 297, Jul. 2024, doi: 10.1145/3657632.
- [8] H. Chen et al., "Conv-Adapter: Exploring Parameter Efficient Transfer Learning for ConvNets," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, Jun. 2024, pp. 1551–1561.
- [9] M. Rastegari, V. Ordonez, J. Redmon, and A. Farhadi, "XNOR-Net: ImageNet Classification Using Binary Convolutional Neural Networks," in *Proceedings of the European Conference on Computer Vision (ECCV)*, Oct. 2016, pp. 525–542.
- [10] I. Hubara et al., "Quantized Neural Networks: Training Neural Networks with Low Precision Weights and Activations," *J. Mach. Learn. Res.*, vol. 18, no. 187, pp. 1–30, Apr. 2018.
- [11] H. Qin et al., "Binary Neural Networks: A Survey," *Pattern Recognit.*, vol. 105, art. no. 107281, Sep. 2020, doi: 10.1016/j.patcog.2020.107281.
- [12] Y. Umuroglu et al., "FINN: A Framework for Fast, Scalable Binarized Neural Network Inference," in *Proceedings of the ACM/SIGDA International Symposium on Field-Programmable Gate Arrays (FPGA)*, Feb. 2017, pp. 65–74.
- [13] M. Blott et al., "FINN-R: An End-to-End Deep-Learning Framework for Fast Exploration of Quantized Neural Networks," *ACM Trans. Reconfigurable Technol. Syst.*, vol. 11, no. 3, art. no. 16, Sep. 2018, doi: 10.1145/3242897.
- [14] S. Jie, H. Wang, and Z.-H. Deng, "Revisiting the Parameter Efficiency of Adapters from the Perspective of Precision Redundancy," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, Oct. 2023, pp. 17217–17226.
- [15] L. Gao, Z. Luo, and L. Wang, "Convolutional Neural Network Acceleration Techniques Based on FPGA Platforms: Principles, Methods, and Challenges," *Information*, vol. 16, no. 10, art. no. 914, Oct. 2025, doi: 10.3390/info16100914.
- [16] A. Nechi et al., "FPGA-Based Deep Learning Inference Accelerators: Where Are We Standing?," *ACM Trans. Reconfigurable Technol. Syst.*, vol. 16, no. 4, art. no. 60, Oct. 2023, doi: 10.1145/3613963.
- [17] S. A. Alam et al., "On the RTL Implementation of FINN Matrix Vector Unit," *ACM Trans. Embed. Comput. Syst.*, vol. 22, no. 6, art. no. 94, Nov. 2023, doi: 10.1145/3547141.
- [18] L. Jungemann, B. Wintermann, H. Riebler, and C. Plessl, "FINN-HPC: Closing the Gap for Energy-Efficient Neural Network Inference on FPGAs in HPC," in *Proceedings of the International Symposium on Highly Efficient Accelerators and Reconfigurable Technologies (HEART)*, May 2025, pp. 103–116.
- [19] J.-F. Schulte et al., "hls4ml: A Flexible, Open-Source Platform for Deep Learning Acceleration on Reconfigurable Hardware," *ACM Trans. Reconfigurable Technol. Syst.*, vol. 19, no. 2, art. no. 19, Jun. 2026, doi: 10.1145/3801979.
- [20] F. Manca, F. Ratto, and F. Palumbo, "ONNX-to-Hardware Design Flow for Adaptive Neural-Network Inference on FPGAs," arXiv:2406.09078, Jun. 2024.
- [21] H.-J. Kang and B.-D. Yang, "AoCStream: All-on-Chip CNN Accelerator with Stream-Based Line-Buffer Architecture and Accelerator-Aware Pruning," *Sensors*, vol. 23, no. 19, art. no. 8104, Sep. 2023, doi: 10.3390/s23198104.
- [22] Z. Zhao et al., "A High-Throughput FPGA Accelerator for Lightweight CNNs With Balanced Dataflow," arXiv:2407.19449, Jul. 2024.
- [23] J. Jiang et al., "FPGA-Based Acceleration for Convolutional Neural Networks: A Comprehensive Review," arXiv:2505.13461, May 2025.
- [24] R. Li, "Dataflow & Tiling Strategies in Edge-AI FPGA Accelerators: A Comprehensive Literature Review," arXiv:2505.08992, May 2025.
- [25] Y. Su, K. P. Seng, L. M. Ang, and J. Smith, "Binary Neural Networks in FPGAs: Architectures, Tool Flows and Hardware Comparisons," *Sensors*, vol. 23, no. 22, art. no. 9254, Nov. 2023, doi: 10.3390/s23229254.
- [26] J.-Y. Zhan et al., "FPGA-Based Acceleration for Binary Neural Networks in Edge Computing," *J. Electron. Sci. Technol.*, vol. 21, no. 2, art. no. 100204, Jun. 2023, doi: 10.1016/j.jnlest.2023.100204.
- [27] P. Antunes and A. Podobas, "FPGA-Based Neural Network Accelerators for Space Applications: A Survey," arXiv:2504.16173, Apr. 2025.
- [28] H.-I. Liu et al., "Lightweight Deep Learning for Resource-Constrained Environments: A Survey," *ACM Comput. Surv.*, vol. 56, no. 10, art. no. 267, Jun. 2024, doi: 10.1145/3657282.
- [29] A. M. Z. Khaki and A. Choi, "Optimizing Deep Learning Acceleration on FPGA for Real-Time and Resource-Efficient Image Classification," *Appl. Sci.*, vol. 15, no. 1, art. no. 422, Jan. 2025, doi: 10.3390/app15010422.
- [30] M. Courbariaux, Y. Bengio, and J.-P. David, "BinaryConnect: Training Deep Neural Networks with Binary Weights During Propagations," in *Proceedings of the Conference on Neural Information Processing Systems (NeurIPS)*, Dec. 2015, pp. 3123–3131.
- [31] Z. Liu et al., "Bi-Real Net: Enhancing the Performance of 1-bit CNNs With Improved Representational Capability and Advanced Training Algorithm," in *Proceedings of the European Conference on Computer Vision (ECCV)*, Sep. 2018, pp. 747–763.
- [32] Z. Liu, Z. Shen, M. Savvides, and K.-T. Cheng, "ReActNet: Towards Precise Binary Neural Network with Generalized Activation Functions," in *Proceedings of the European Conference on Computer Vision (ECCV)*, Aug. 2020, pp. 143–159.
- [33] S. Liang et al., "FP-BNN: Binarized Neural Network on FPGA," *Neurocomputing*, vol. 275, pp. 1072–1086, Jan. 2018, doi: 10.1016/j.neucom.2017.09.046.
- [34] B. Jacob et al., "Quantization and Training of Neural Networks for Efficient Integer-Arithmetic-Only Inference," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, Jun. 2018, pp. 2704–2713.
- [35] M. Nagel et al., "A White Paper on Neural Network Quantization," arXiv:2106.08295, Jun. 2021.
- [36] L. Zhang et al., "Quantized Convolutional Neural Networks: A Hardware Perspective," *Front. Electron.*, vol. 6, art. no. 1469802, Jul. 2025, doi: 10.3389/felec.2025.1469802.
- [37] M. Tasci, A. Istanbulu, V. Tumen, and S. Kosunalp, "FPGA-QNN: Quantized Neural Network Hardware Acceleration on FPGAs," *Appl. Sci.*, vol. 15, no. 2, art. no. 688, Jan. 2025, doi: 10.3390/app15020688.
- [38] C. Ding et al., "LoRA-C: Parameter-Efficient Fine-Tuning of Robust CNN for IoT Devices," arXiv:2410.16954, Oct. 2024.
- [39] H. Kwak, K. Lee, J.-J. Lee, and W. Lee, "LoRA-Edge: Tensor-Train-Assisted LoRA for Practical CNN Fine-Tuning on Edge Devices," arXiv:2511.03765, Nov. 2025.
- [40] T. Lu and C. Chang, "Cache-Adapter: Efficient Video Action Recognition Using Adapter Fine-Tuning and Cache Memorization Technique," *J. Vis. Commun. Image Represent.*, vol. 111, art. no. 104543, Sep. 2025, doi: 10.1016/j.jvcir.2025.104543.
- [41] R. Belanec, B. Pecher, I. Srba, and M. Bielikova, "PEFT-Bench: A Parameter-Efficient Fine-Tuning Methods Benchmark," arXiv:2511.21285, Nov. 2025.
- [42] A. Boudjadar, S. Islam, and R. Buyya, "Dynamic FPGA Reconfiguration for Scalable Embedded AI: A Co-Design Methodology for CNN Acceleration," *Future Gener. Comput. Syst.*, vol. 169, art. no. 107777, Aug. 2025, doi: 10.1016/j.future.2025.107777.
- [43] M. Kang and D. Park, "Runtime-Robust Edge Inference System with Masking-Based Partial Update on Dynamic Reconfigurable FPGA," *Sensors*, vol. 25, no. 24, art. no. 7448, Dec. 2025, doi: 10.3390/s25247448.
- [44] A. Pappalardo, "Brevitas: A Quantization-Aware Training Library," 2021. [Online]. Available: <https://github.com/Xilinx/brevitas>
- [45] B. A. Arend, A. F. Ely, F. L. Kastensmidt, and M. Recamonde-Mendoza, "Predicting Hardware Resource Allocation of MVAU Modules in FINN-Generated FPGA Accelerators," in *Proceedings of the I Escola Regional de Aprendizagem de Máquina e Inteligência Artificial da Região Sul (ERAMIA-RS)*, Nov. 2025, pp. 308–311.
- [46] K. He, X. Zhang, S. Ren, and J. Sun, "Deep Residual Learning for Image Recognition," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, Jun. 2016, pp. 770–778.