

MARS: A Reconfigurable Multi-Task-Capable Accelerator Architecture for Ultra-Low-Bit Streaming Inference

Po-Chun Huang*, Pao-Ann Hsiung*, and Chun-Hsian Huang[†]

*Department of Computer Science and Information Engineering, National Chung Cheng University, Chiayi, Taiwan

[†]Department of Electrical Engineering, National Chung Cheng University, Chiayi, Taiwan

Abstract—Common approaches to serving multiple AI tasks on one edge FPGA rely either on partial reconfiguration—with a fabric controller, per-task bitstreams, and millisecond-to-hundred-millisecond switches—or task sets frozen at synthesis time. We present MARS, a 1-bit-weight/1-bit-activation streaming accelerator in which one frozen backbone is shared by all tasks and a task switch is only a ≈ 25 KB parameter rewrite. Two algorithmic pieces make the ultra-low-bit adapter path work: a Residual Correction (RC) term (the adapter’s pre-existing down-projection bias), with a substantially larger measured effect at 1 bit (up to +23.0 pp) and realised at zero additional arithmetic or datapath cost by folding it into the accumulator initial value, and multi-branch adapter scaling, which increases the effective capacity of the binary adapter path (+3.4 pp mean, up to +11.3 pp, across twenty backbone–target pairs in the deployed-geometry sweep). A 19-LUT Function Switch Controller reprograms all task-dependent parameters over one bitstream. On a PYNQ-Z2, five task-specific models switch at runtime with on-board accuracy within 0.69 pp of software and a 1.86 ms configuration rewrite with verified post-switch correctness; a throughput build sustains 1,866 img/s at 2.214 W (842.7 img/s/W) on a protocol-matched 1W1A workload with deployment-representative, different adapter geometries (Vivado-estimated FPGA power vs. measured GPU power).

Index Terms—FPGA accelerator, ultra-low-bit quantisation, parameter-efficient transfer learning, runtime task switching

I. INTRODUCTION

Edge FPGAs increasingly serve several AI tasks from one device. The prevailing mechanism, partial reconfiguration (PR), swaps hardware regions per task and pays a fixed tax: a reconfiguration controller, a per-task partial-bitstream repository, and measured switch latencies of 5.2 ms (ZyCAP) [1] to 97–114 ms (PCAP) [2] on PYNQ-class devices. Fixed-feature-extractor alternatives [3] and on-chip variant selection [4] avoid fabric rewrites but freeze the reachable task set at synthesis.

Parameter-efficient transfer learning (PEFT) [5], [6] suggests a third path: freeze one backbone and adapt tasks through small modules such as Conv-Adapter [7]. If the backbone is binarised (1W1A) [8], [9] and mapped to a FINN-style streaming pipeline [10], a task switch reduces to rewriting a small parameter set—no fabric reconfiguration at all. We find, however, that at 1 bit a single binary adapter branch exhibits a pronounced transfer-accuracy degradation, absent at comparable magnitude at ≥ 2 bits [11].

MARS (Multi-Branch Adapter Reconfigurable Streaming-accelerator) makes the ultra-low-bit adapter path practical and verifies it on hardware. Contributions: (i) the 1-bit-specific Residual Correction finding and its zero-additional-arithmetic hardware realisation; (ii) multi-branch adapter scaling for binary networks; (iii) single-bitstream runtime task switching for eligible same-backbone tasks at $O(1)$ on-chip cost via a 19-LUT controller; and (iv) on-board verification with explicit evidence boundaries—five tasks within 0.69 pp of software, 1.86 ± 0.04 ms configuration-rewrite switching (post-switch correctness verified), while $M \geq 2$ branch configurations exceed the XC7Z020 and are verified at the RTL-simulation and post-synthesis levels only. To the best of our knowledge among the surveyed FPGA/PEFT accelerator works, MARS is the first to combine an ultra-low-bit shared backbone, feature-extractor-level adapters, and single-bitstream runtime task switching verified on hardware.

II. RELATED WORK

Ultra-low-bit dataflow. BNNs replace multiply–accumulate with XNOR–popcount [12], [8], [9], with representational improvements in [13], [14] and FPGA systems such as FP-BNN [15], and FINN-style compilers [10], [16], [17] and related flows [18] give each layer a dedicated compute unit with on-chip weights; surveys catalogue this design space [19], [20], [21], [22]—but compile one task per bitstream. **PEFT.** Adapters [5], LoRA [6], and Conv-Adapter [7] adapt frozen backbones with few parameters [23], with edge-oriented CNN variants emerging [24], [25], yet target full or moderate precision; at 1 bit the adapter itself must be binary, and its transfer effectiveness can degrade sharply [11]. **Task switching on FPGAs.** PR systems [2], [1] pay controller, storage, and latency costs; fixed-feature-extractor [3] and variant-selection [4] designs freeze the task set at synthesis. MARS occupies the unfilled point: an open set of eligible same-backbone tasks, no fabric reconfiguration, and an ultra-low-bit shared backbone.

III. THE MARS METHOD

A. Binary Adapter with Residual Correction

MARS inserts one adapter beside each of the five adapter-hosting convolution layers of a frozen 1W1A VGG-like backbone (FINN’s CNV: six conv + three FC layers, pre-trained on




Fig. 1. RC in hardware: accumulator initialisation from `ram_rc`; no extra adder or pipeline stage.

CIFAR-10). The down-projection accumulates XNOR products with 1-bit weights plus a per-channel bias,

$$h_j = b_j + \sum_{i=1}^{C_{in}} \text{XNOR}(W_{j,i}^{\text{down}}, x_i), \quad (1)$$

followed by a binary activation $a_j = \mathbb{I}[h_j \geq 0]$ and a 1-bit up-projection $z_o = \sum_j \text{XNOR}(W_{o,j}^{\text{up}}, a_j)$. The bias b_j is the Residual Correction: at 1 bit, without the bias, the binary activation uses a fixed channel-wise cut that may be misaligned with the task-dependent activation distribution; RC shifts this cut. This interpretation is consistent with the observed ablation, although the complete causal pathway to final accuracy is not isolated here. In hardware, RC becomes the *initial value* of the down-projection accumulator—a load the accumulator performs every window anyway—so it costs zero additional arithmetic, DSPs, or pipeline stages; only its storage and configuration bytes count, and they are included in the adapter budget. Fig. 1 shows the realisation: `ram_rc` drives the initialisation input of the accumulator multiplexer, selected at every window start.

B. Multi-Branch Scaling and Training

MARS sums M structurally identical branches with independent weights, RC, and scales,

$$y_o = \sum_{i=1}^M \alpha_i z_o^{(i)}, \quad (2)$$

enriching the integer-valued corrections the adapter can express (each branch output is itself multi-level; no fixed level count is claimed). The overhead is linear: +58,533 parameters (+3.78% of the backbone) and +4.23 MMACs per branch (compute-then-crop software reference; the deployed RTL’s centre-selection evaluates 3.01 MMACs). Training freezes the backbone and updates the adapter and classifier jointly: weights binarise in the forward pass ($W = \text{sign}(\tilde{W})$) with the straight-through estimator $\partial \text{sign}(u)/\partial u \approx \mathbf{1}_{|u| \leq 1}$ in the backward pass, while each real-valued scale receives the exact gradient $\partial L/\partial \alpha_i = \langle \partial L/\partial y, z^{(i)} \rangle$. At export, α_i is quantised into a Q8 contribution look-up table and RC into the accumulator initial value.

IV. HARDWARE ARCHITECTURE

A. Streaming Pipeline and Super Wrapper

Fig. 2 shows the deployed system on a Zynq XC7Z020 at 100 MHz: MVAU0, five adapter-enhanced Super Wrappers




Fig. 2. MARS streaming dataflow: MVAU0, adapter-enhanced Super1–5, MVAU6–8, and the Function Switch Controller.




Fig. 3. Super Wrapper: the modified MVAU core exposes raw partial sums; the adapter contribution is fused before the shared Q8 threshold.

(Super1–5), and MVAU6–8, with a Function Switch Controller (FSC) fanning a configuration bus out to all task-dependent memories. In a FINN pipeline, thresholding happens inside each matrix–vector–activation unit, so raw partial sums are not normally exposed; the Super Wrapper (Fig. 3) decouples that thresholding, duplicates the input stream to the modified MVAU core and the Adapter IP, synchronises the adapter side-stream through a small FIFO, and fuses both before one shared Q8 threshold:

$$\hat{y}_o = \mathbb{I}[(p_o \ll 8) + s_o \cdot \text{LUT}_\alpha(q_o) \geq T_o^{Q8}], \quad (3)$$

with p_o the backbone partial sum, q_o the up-projection popcount, and T_o^{Q8} the BatchNorm-fused threshold—shifts, look-ups, and integer adds only. All adapter memories are configuration-writable distributed RAM, so the adapter adds zero BRAM and zero DSPs.

B. Runtime Task Switching

The FSC is a 19-LUT AXI-Lite slave that decodes a 64 KB aperture into nine destinations (MVAU0 thresholds, MVAU8 classifier weights, five adapter blobs, MVAU6/7 thresholds). Because all eligible tasks share the frozen backbone, only BatchNorm-fused thresholds, the classifier projection, and adapter parameters differ per task: a full switch is a ≈ 25 KB rewrite (25,088 stored B; 6,757 32-bit writes = 27,028 bus B, sub-word fields widened on the bus)—wait for the output-DMA done flag, bulk-write the blob, launch the next batch; no additional flush beyond the output-DMA-done barrier, no fabric reconfiguration (Fig. 4). Each additional task is an off-line host-side blob: on-chip cost is $O(1)$ in the task count for same-backbone, $32 \times 32 \times 3$, same-output-head classification tasks. Two builds realise two depths of this mechanism: a *throughput* build (FINN-default folding, one adapter task baked at synthesis, 2-task) and a *compactness* build (uniform




Fig. 4. Runtime task-switching sequence.

TABLE I
RC ABLATION AT 1 BIT, WIDE GEOMETRY (CIFAR-10 $\rightarrow$ SVHN).

Adapter	M	No-RC	With RC	Δ_{RC}	Full-FT
Conv-Adapter [7]	1	63.06%	79.23%	+16.17	
MARS	2	66.12%	83.01%	+16.89	94.91%
	3	67.79%	83.48%	+15.69	
	4	69.08%	84.44%	+15.36	

PE=1, memories power up empty, any eligible same-backbone task loadable at runtime, N -task).

V. EXPERIMENTAL RESULTS

Setup. PyTorch + Brevitas [26] QAT (200 epochs, ADAM, lr 0.005, seed 2024); FINN + Vivado 2022.2; PYNQ-Z2 (XC7Z020). Five $32 \times 32 \times 3$ datasets: CIFAR-10 (backbone), SVHN, FashionMNIST, STL10, CINIC10. Correctness is verified at three levels: per-module RTL simulation (43,520 vectors, 100% bit-exact), top-level dataflow simulation against the PyTorch golden model, and on-board execution. Accuracies are best-epoch operating points selected on the test split (original BNN-PYNQ trainer flow; the identical rule for every configuration).

RC at 1 bit. Table I isolates RC on CIFAR-10 $\rightarrow$ SVHN: +15 to +17 pp at every branch count in the wide geometry and +23.0 pp in the deployed geometry, while at ≥ 2 bits the same term is typically within ± 0.3 pp (maximum observed 1.4 pp). Gains on other targets range from +0.6 to +7.4 pp.

Multi-branch scaling. Across twenty backbone-target pairs spanning five source backbones, the deployed-geometry sweep (Table II) raises 1-bit transfer accuracy by +3.36 pp on average over the single-branch baseline, up to +11.33 pp (FashionMNIST $\rightarrow$ SVHN), with the largest gains on hard transfers and diminishing returns as M grows. Table III shows the complementary wide-geometry software study on the CIFAR-10 backbone. Every one of the twenty pairs improves (largest: FashionMNIST $\rightarrow$ SVHN +11.33 pp, STL10 $\rightarrow$ SVHN +7.50 pp, CINIC10 $\rightarrow$ SVHN +7.20 pp); five paired seeds provide preliminary evidence of a positive $M=1 \rightarrow 4$ gain on two deployment-relevant transfers under the same test-selected checkpoint protocol: $+6.64 \pm 0.43$ pp on SVHN and $+2.69 \pm 0.36$ pp on FashionMNIST; a selection-free final-epoch reading of the same logs yields consistent gains ($+6.87/ +2.99$ pp).

On-board task switching. Table IV: the five task-specific models served from the single compactness bitstream (loaded in turn from host-resident images) each reproduce their software reference within 0.69 pp (10,000-image test sets; STL10 its full 8,000); the residual gap is consistent with fixed-point

TABLE II
DEPLOYED-GEOMETRY MULTI-BRANCH TRANSFER ACCURACY (%) OVER ALL TWENTY BACKBONE-TARGET PAIRS ($M=1$ vs. $M=4$, RC ON; SINGLE SEED). $\Delta_{MULTI} = M=4 - M=1$ (pp).

Backbone	Target	$M=1$	$M=4$	Δ_{multi}
CIFAR-10	SVHN	73.72	79.81	+6.09
	FashionMNIST	80.42	83.54	+3.12
	STL10	67.46	68.06	+0.60
	CINIC10	64.80	65.45	+0.65
SVHN	CIFAR-10	43.18	48.37	+5.19
	FashionMNIST	78.51	81.89	+3.38
	STL10	34.41	35.10	+0.69
	CINIC10	33.59	38.03	+4.44
STL10	CIFAR-10	66.65	68.79	+2.14
	FashionMNIST	81.48	84.19	+2.71
	SVHN	68.15	75.65	+7.50
	CINIC10	54.26	56.56	+2.30
FashionMNIST	CIFAR-10	33.63	36.76	+3.13
	STL10	32.64	34.11	+1.47
	SVHN	51.37	62.70	+11.33
	CINIC10	27.75	29.56	+1.81
CINIC10	CIFAR-10	78.21	78.72	+0.51
	SVHN	69.14	76.34	+7.20
	STL10	67.21	67.42	+0.21
	FashionMNIST	80.23	82.88	+2.65
Mean gain over $M=1$ (pp)		—	—	+3.36

TABLE III
MULTI-BRANCH SCALING ON THE CIFAR-10 BACKBONE (WIDE GEOMETRY, SOFTWARE, 1W1A, RC).

Target	[7]		MARS				Full-FT
	$M=1$	$M=2$	$M=3$	$M=4$			
SVHN	79.23	83.01	83.48	84.44			94.91
STL10	67.50	67.48	67.68	68.15			71.05
FashionMNIST	82.71	83.57	83.63	83.60			92.36
CINIC10	64.50	64.71	64.85	65.02			68.02




Fig. 5. Accuracy gain vs. post-synthesis LUT demand for branch scaling (compactness build; large-gain: $\Delta_{multi} \geq 3$ pp, 9 pairs vs. 11); $M \geq 2$ exceeds the XC7Z020.

TABLE IV
ON-BOARD N -TASK SWITCHING ACCURACY (COMPACTNESS BUILD).

Dataset	Mode	Software Reference	On-Board	Gap (pp)
CIFAR-10	backbone only	81.46%	80.99%	-0.47
SVHN	MARS	73.72%	73.21%	-0.51
FashionMNIST	MARS	80.42%	79.98%	-0.44
STL10	MARS	67.46%	66.77%	-0.69
CINIC10	MARS	64.80%	64.64%	-0.16

TABLE V
TASK-SWITCHING FPGA INFERENCE VS. PUBLISHED APPROACHES.
LATENCIES CORRESPOND TO DIFFERENT SWITCHING SEMANTICS; THE
ADAPTATION LEVEL AND RESIDENCY MODEL DIFFER ACROSS
APPROACHES.

Approach	Switch Latency	PR Ctrl.	Task Storage Scaling
ACNNE [2]	97–114 ms	Yes	partial bitstream / task
CoarseToFine [1]	5.2 ms	Yes	partial bitstream / task
FixyNN [3]	N/A	No	back-end HW / task
NeuroAdaptive [4]	0.23–0.67 μ s	No	on-chip, fixed at synthesis
MARS	1.86 ms	No	$O(1)$ on-chip; 25 KB blob / task

TABLE VI
PER-BRANCH PARAMETER AND POST-SYNTHESIS COST (COMPACTNESS).

M	Params	vs. Backbone	Slice LUTs (%)	Power (W)
1	58,533	+3.78%	47,931 (90.1%)	1.841
2	117,066	+7.57%	71,841 (135.0%)	1.897
3	175,599	+11.35%	95,593 (179.7%)	1.915
4	234,132	+15.14%	125,538 (236.0%)	1.969

TABLE VII
TWO-TASK VS. N -TASK SWITCHING.

	Throughput (2-task)	Compactness (N -task)
Folding	heterog. PE ≤ 32	uniform PE = 1
Baked at synthesis	one adapter task	nothing
Moved per switch	≈ 15 KB	full blob ≈ 25 KB
Reachable models	exactly two	any eligible same-backbone task
Verified	on-board, 2 tasks	on-board, 5 tasks

and threshold-quantisation effects; no mismatch was observed in the per-module RTL tests. Fifty random switches give a 1.86 ± 0.04 ms configuration rewrite (post-switch correctness verified separately), $\approx 0.02\%$ of a batch-1000 run. Compared with published task-switching FPGA inference (Table V), MARS is the only surveyed system needing no fabric-PR controller while keeping the task set open after synthesis.

Table VI itemises the per-branch cost, and Table VII contrasts how the two builds realise two depths of the same switching mechanism.

Fig. 7 shows the post-implementation placement of the compactness build: adapter logic and backbone MVAUs dominate and interleave across the die, while the FSC appears only as a negligible scattering of cells, matching its 19-LUT footprint. The 1.86 ms switch amortises to $\approx 0.02\%$ of a batch-1000 run; serving T tasks requires one bitstream plus T host-side 25 KB blobs, versus T partial bitstreams for PR systems.




Fig. 6. Folding difference between the builds: PE=1 (compactness) vs. PE-parallel lanes (throughput, illustrated for PE=4).




Fig. 7. Layout of the compactness build (XC7Z020). Yellow: adapter + adapter RAM; green: MVAUs; purple: FSC; orange: PS; blue: other SoC logic.

TABLE VIII
DEPLOYMENT-LEVEL COMPARISON, MATCHED PROTOCOL BUT DIFFERENT
ADAPTER GEOMETRIES (SVHN, BATCH 1000).

Metric	RTX 4090	Orin NX	MARS (Compact.)	MARS (Thruput.)
Adapter geometry	3×3 wide	3×3 wide	1×1 depl.	1×1 depl.
FPS (img/s)	9,388	2,752	107.1	1,866
Power (W)	132.48	19.23	1.802	2.214
Efficiency (img/s/W)	70.86	143.1	59.4	842.7

Cross-platform. On a protocol-matched 1W1A SVHN workload at batch 1000 (Table VIII), the throughput build sustains 1,866 img/s at 2.214 W—842.7 img/s/W, 11.9× the RTX 4090 and 5.9× the Jetson Orin NX. FPGA power is the Vivado post-route estimate while GPU/Jetson power is runtime-measured (NVML/tegrastats), so the ratios are specific to this workload and measurement methodology, not a general platform-efficiency claim; the GPU/Jetson run the 3×3 Conv-Adapter geometry while the FPGA runs the deployed 1×1 MARS geometry (a deployment-level comparison, not a same-model sweep).

VI. DISCUSSION AND LIMITATIONS

The two builds are complementary operating points: the throughput build for a fixed task (or task pair) known at synthesis, the compactness build when the task set must stay

open—its throughput cost is the PE=1 folding (Fig. 6), not the switching machinery, and the adapter-enabled throughput build stays within $\approx 0.5\%$ of its backbone because the adapter side-stream is cycle-accurately synchronised into the folded datapath. Distributed-RAM adapter memories keep BRAM for the backbone memstreams; the selected eight-fractional-bit (Q8) scale keeps the measured board-to-software gap below one point in the deployed experiments; a systematic precision sweep was not performed.

The evidence boundaries are explicit. The hardware-verified point is the $M=1$ compactness build; $M \geq 2$ exceeds the XC7Z020 and is RTL-simulation/post-synthesis-verified only. The 1.86 ms switch is bound by single-beat AXI-Lite writes—a faster bulk driver corrupted adapter state, so sub-millisecond switching requires an RTL change. “Arbitrary” task counts refer to the resource dimension only: eligible tasks share the frozen backbone, the $32 \times 32 \times 3$ input, the output-head interface, and the classification setting; host-side blob storage grows linearly; accuracy remains bounded by the shared backbone’s transfer capacity. FPGA power is a Vivado estimate rather than a board measurement.

VII. CONCLUSION

MARS combines ultra-low-bit streaming inference with parameter-efficient task adaptation on a single FPGA bitstream: a 1-bit-specific residual correction realised at zero additional arithmetic cost, multi-branch adapters that increase the effective capacity of the binary adapter path, and a 19-LUT controller that turns task switching into a 25 KB parameter rewrite. Five task-specific models run on one PYNQ-Z2 compactness-build bitstream within 0.69 pp of software, with 1.86 ms configuration-rewrite switching (post-switch correctness verified); separately, the throughput build reaches 842.7 img/s/W on the matched workload. Limitations are explicit: $M \geq 2$ is verified at the RTL-simulation and post-synthesis levels only (not on-board), switching is not yet sub-millisecond, eligible tasks share the frozen backbone, input format, and output-head interface, accuracies are best-epoch test-selected operating points, and FPGA power is a Vivado estimate. Future work targets a burst-capable switching driver, multi-branch verification on larger devices, and residual backbones.

REFERENCES

- [1] A. Rangasikunpam, S. Amiri, and L. Ost, “A Reconfigurable Coarse-to-Fine Approach for the Execution of CNN Inference Models in Low-Power Edge Devices,” *IET Comput. Digit. Tech.*, vol. 2024, no. 1, art. no. 6214436, Jan. 2024, doi: 10.1049/cdt2/6214436.
- [2] C.-H. Huang, S.-W. Tang, and P.-A. Hsiung, “ACNNE: An Adaptive Convolution Engine for CNNs Acceleration Exploiting Partial Reconfiguration on FPGAs,” in *Proceedings of the IEEE International Symposium on Circuits and Systems (ISCAS)*, May 2024, pp. 1–5.
- [3] P. N. Whatmough et al., “FixyNN: Energy-Efficient Real-Time Mobile Computer Vision Hardware Acceleration via Transfer Learning,” in *Proceedings of the Conference on Machine Learning and Systems (MLSys)*, Apr. 2019.
- [4] A. El Bouazzaoui, O. Mouhib, and A. Hadjoudja, “NeuroAdaptiveNet: A Reconfigurable FPGA-Based Neural Network System with Dynamic Model Selection,” *Chips*, vol. 4, no. 2, art. no. 24, May 2025, doi: 10.3390/chips4020024.
- [5] N. Houlsby et al., “Parameter-Efficient Transfer Learning for NLP,” in *Proceedings of the International Conference on Machine Learning (ICML)*, Jun. 2019, pp. 2790–2799.
- [6] E. J. Hu et al., “LoRA: Low-Rank Adaptation of Large Language Models,” in *Proceedings of the International Conference on Learning Representations (ICLR)*, Apr. 2022.
- [7] H. Chen et al., “Conv-Adapter: Exploring Parameter Efficient Transfer Learning for ConvNets,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, Jun. 2024, pp. 1551–1561.
- [8] M. Rastegari, V. Ordonez, J. Redmon, and A. Farhadi, “XNOR-Net: ImageNet Classification Using Binary Convolutional Neural Networks,” in *Proceedings of the European Conference on Computer Vision (ECCV)*, Oct. 2016, pp. 525–542.
- [9] I. Hubara et al., “Quantized Neural Networks: Training Neural Networks with Low Precision Weights and Activations,” *J. Mach. Learn. Res.*, vol. 18, no. 187, pp. 1–30, Apr. 2018.
- [10] Y. Umuroglu et al., “FINN: A Framework for Fast, Scalable Binarized Neural Network Inference,” in *Proceedings of the ACM/SIGDA International Symposium on Field-Programmable Gate Arrays (FPGA)*, Feb. 2017, pp. 65–74.
- [11] S. Jie, H. Wang, and Z.-H. Deng, “Revisiting the Parameter Efficiency of Adapters from the Perspective of Precision Redundancy,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, Oct. 2023, pp. 17217–17226.
- [12] M. Courbariaux, Y. Bengio, and J.-P. David, “BinaryConnect: Training Deep Neural Networks with Binary Weights During Propagations,” in *Proceedings of the Conference on Neural Information Processing Systems (NeurIPS)*, Dec. 2015, pp. 3123–3131.
- [13] Z. Liu et al., “Bi-Real Net: Enhancing the Performance of 1-bit CNNs With Improved Representational Capability and Advanced Training Algorithm,” in *Proceedings of the European Conference on Computer Vision (ECCV)*, Sep. 2018, pp. 747–763.
- [14] Z. Liu, Z. Shen, M. Savvides, and K.-T. Cheng, “ReActNet: Towards Precise Binary Neural Network with Generalized Activation Functions,” in *Proceedings of the European Conference on Computer Vision (ECCV)*, Aug. 2020, pp. 143–159.
- [15] S. Liang et al., “FP-BNN: Binarized Neural Network on FPGA,” *Neurocomputing*, vol. 275, pp. 1072–1086, Jan. 2018, doi: 10.1016/j.neucom.2017.09.046.
- [16] M. Blott et al., “FINN-R: An End-to-End Deep-Learning Framework for Fast Exploration of Quantized Neural Networks,” *ACM Trans. Reconfigurable Technol. Syst.*, vol. 11, no. 3, art. no. 16, Sep. 2018, doi: 10.1145/3242897.
- [17] S. A. Alam et al., “On the RTL Implementation of FINN Matrix Vector Unit,” *ACM Trans. Embed. Comput. Syst.*, vol. 22, no. 6, art. no. 94, Nov. 2023, doi: 10.1145/3547141.
- [18] J.-F. Schulte et al., “hls4ml: A Flexible, Open-Source Platform for Deep Learning Acceleration on Reconfigurable Hardware,” *ACM Trans. Reconfigurable Technol. Syst.*, vol. 19, no. 2, art. no. 19, Jun. 2026, doi: 10.1145/3801979.
- [19] H. Qin et al., “Binary Neural Networks: A Survey,” *Pattern Recognit.*, vol. 105, art. no. 107281, Sep. 2020, doi: 10.1016/j.patcog.2020.107281.
- [20] Y. Su, K. P. Seng, L. M. Ang, and J. Smith, “Binary Neural Networks in FPGAs: Architectures, Tool Flows and Hardware Comparisons,” *Sensors*, vol. 23, no. 22, art. no. 9254, Nov. 2023, doi: 10.3390/s23229254.
- [21] J. Jiang et al., “FPGA-Based Acceleration for Convolutional Neural Networks: A Comprehensive Review,” arXiv:2505.13461, May 2025.
- [22] H.-I. Liu et al., “Lightweight Deep Learning for Resource-Constrained Environments: A Survey,” *ACM Comput. Surv.*, vol. 56, no. 10, art. no. 267, Jun. 2024, doi: 10.1145/3657282.
- [23] R. Belanec, B. Pecher, I. Srba, and M. Bielikova, “PEFT-Bench: A Parameter-Efficient Fine-Tuning Methods Benchmark,” arXiv:2511.21285, Nov. 2025.
- [24] C. Ding et al., “LoRA-C: Parameter-Efficient Fine-Tuning of Robust CNN for IoT Devices,” arXiv:2410.16954, Oct. 2024.
- [25] H. Kwak, K. Lee, J.-J. Lee, and W. Lee, “LoRA-Edge: Tensor-Train-Assisted LoRA for Practical CNN Fine-Tuning on Edge Devices,” arXiv:2511.03765, Nov. 2025.
- [26] A. Pappalardo, “Brevitas: A Quantization-Aware Training Library,” 2021. [Online]. Available: <https://github.com/Xilinx/brevitas>